

Clustering

Rui

May 22, 2023

目录

1	K-Means	2
1.1	Computation	2
1.2	Visualization	3
2	Hierachical Clustering	5
3	No. of Clusters	9
3.1	Within-Cluster SS	9
3.2	Silhouette Method	10
3.3	Majority Rule	11
4	Method Agreement Procedure:	12
5	Two Step Clustering	12
6	Clustering for Cases and Variables	13
6.1	Q1	13
6.2	Q2	15

```
library("parameters") #for clustering summary
library("see")
library("factoextra") #for clustering plotting
library(bruceR)
```

```
library(tidyverse)
```

```
athlete = import("data13-02.xlsx") %>% select(-1)
```

```
## Successfully imported: 10 obs. of 4 variables
```

1 K-Means

1.1 Computation

- **x**: numeric matrix-like data
- **centers**: either the number of clusters, say `kk`, or a set of initial (distinct) cluster centres. If a number, a random set of (distinct) rows in `x` is chosen as the initial centres.
- **iter.max**: the maximum number of iterations allowed.
- **nstart**: if `centers` is a number, how many random sets should be chosen. Trying several random starts is often recommended.

```
kmean_results = kmeans(athlete, centers = 4,  
                        iter.max = 10, nstart = 1)
```

`kmeans` returns an object of class “`kmeans`” which has a `print` and a `fitted` method. It is a list with at least the following components:

- **cluster** : A vector of integers (from 1:k) indicating the cluster to which each point is allocated.
- **centers**: A matrix of cluster centres.
- **totss**: The total sum of squares.
- **withinss**: Vector of within-cluster sum of squares, one component per cluster.
- **tot.withinss**: Total within-cluster sum of squares, i.e. `sum(withinss)`.
- **betweenss**: The between-cluster sum of squares, i.e. `totss - tot.withinss`.
- **size**: The number of points in each cluster.

```
kmean_results$cluster #group labels for each observation
```

```
## [1] 3 1 4 3 1 2 4 1 4 2
```

```
kmean_results$centers #centers
```

```
##          x1          x2          x3
## 1 121.6667 18.33333 43.00000
## 2 120.5000 19.00000 44.50000
## 3 124.5000 20.00000 44.50000
## 4 121.0000 17.00000 41.66667
```

```
kmean_results$totss # total variance
```

```
kmean_results$withinss # within-group variance
```

```
kmean_results$tot.withinss # sum of within SS
```

```
kmean_results$betweenss #between-group variance
```

```
## [1] 51.6
```

```
## [1] 1.333333 1.000000 1.000000 2.666667
```

```
## [1] 6
```

```
## [1] 45.6
```

```
kmean_results$size # how many cases in each cluster
```

```
## [1] 3 2 2 3
```

1.2 Visualization

Use `factoextra::fviz_cluster(object, data)`. For more details about useful parameters, see the help document. Remember that `object` and `data` are two most indispensable parameters.

```
fviz_cluster(kmean_results, data = athlete)
```

```
## Registered S3 methods overwritten by 'broom':
```

```
##   method          from
```

```
## tidy.glm         jtools
```

```
## tidy.summary.glht jtools
```




One way to start is from random centers; another way is to specify a few sets of centers and start from them.

Clearly, starting from different centers, either random or specified, the results may vary, though not significant. If the iteration will converge eventually, then given sufficiently-large iteration time, the center will be determined and fixed.

```
initial_centers = import("data13-02a.xlsx") %>% select(-1)
```

```
## Successfully imported: 4 obs. of 4 variables
```

```
kmean_results_initialized = kmeans(athlete, initial_centers ,
                                   iter.max = 10, nstart = 1)
fviz_cluster(kmean_results_initialized, data = athlete)
```



2 Hierarchical Clustering

```
automobile = import("data13-01.xlsx") %>%
  select(c('type', 'price', 'engine_s', 'horsepow',
           'wheelbas', 'width', 'length', 'curb_wgt',
           'fuel_cap', 'mpg')) %>% drop_na() %>% scale()
```

Successfully imported: 157 obs. of 27 variables


Use “`hclust()`” from base R and transmit a **dissimilarity** structure produced by “`dist()`”.

```
distMat = dist(automobile, method = "euclidean") # Euclidean distance matrix
hc = hclust(distMat) # hierarchical clustering, using the default method
```

Use “`plot(hclust_obj)`” to visualize the hierarchical clustering. Note that you should specify “`hang = -1`” to ensure all nodes touch the same horizontal line and looks decent. If the sample isn’t that large, set “`labels =`


TRUE” (by default) to see the corresponding observations.

```
plot(hc, labels=FALSE, hang = -1)
```



```
plot(hc, labels=FALSE)
```


Cluster Dendrogram



```
distMat
hclust (*, "complete")
```

Moreover, you can draw dendrogram with grouping.

```
fviz_dend(hc,k=3, rect = T)
```



Also, you can also check how many members for each cluster and make a frequency table.

Note that the grouping might change if we change cluster number. And this in essence is the pros of hierarchical clustering.

```
memberLabels = cutree(hc,3)
table(memberLabels)
```

```
## memberLabels
##  1  2  3
## 71 71 10
```

Then, we can get the centers of clusters by “mean” summary statistics. Use “aggregate()” to achieve that.

- x: usually a data frame.
- by: a list of grouping elements, each **as long** as the variables in the data frame x. The elements are coerced to factors before use.
- FUN: a function to compute the summary statistics which can be applied to all data subsets.

```
aggregate(automobile,list(memberLabels),mean)
```


```
##  Group.1      type      price  engine_s  horsepower  wheelbas      width
## 1      1 -0.4362472 -0.4985325 -0.7231739 -0.6579479 -0.6259158 -0.6459818
## 2      2  0.5201409  0.1193177  0.5017194  0.3293009  0.6601116  0.5979645
## 3      3 -0.5956452  2.6924245  1.5723270  2.3333938 -0.2427902  0.3409229
##      length  curb_wgt  fuel_cap      mpg
## 1 -0.65105171 -0.7853229 -0.7200528  0.6620751
## 2  0.65553921  0.7236353  0.6270241 -0.5592928
## 3 -0.03186123  0.4379815  0.6605042 -0.7297543
```

Moreover, we can use Silhouette plot to check whether some cases are outliers.

```
library(cluster)
plot(silhouette(cutree(hc,3), distMat))
```


Silhouette plot of (x = cutree(hc, 3), dist = distMat)

n = 152


3 clusters C_j $j : n_j \mid \text{ave}_{i \in C_j} s_i$ 

3 No. of Clusters

3.1 Within-Cluster SS

Looking for the “elbow” where the SS starts to slow down its drop.

```
fviz_nbclust(automobile, kmeans, method = "wss", print.summary=TRUE)
```




here suggests 2 or 3.

3.2 Silhouette Method

Looking for the value that sets the highest average silhouette width, which indicates the optimal number of clusters.

```
fviz_nbclust(automobile, kmeans, method = "silhouette", print.summary=TRUE)
```



3.3 Majority Rule

`n_clusters()` is the main function to find out the optimal numbers of clusters present in the data based on the maximum consensus of a large number of methods.

Essentially, there exist many methods to determine the optimal number of clusters, each with pros and cons, benefits and limitations. The main `n_clusters()` function proposes to run all of them, and find out the number of clusters that is suggested by the **majority** of methods (in case of ties, it will select the most parsimonious solution with fewer clusters).

Note that we also implement some specific, commonly used methods, like the Elbow or the Gap method, with their own visualization functionalities.


Also note that the `n_clusters()` function will do the normalization automatically (by default).

```
n_clusters(data.frame(automobile))
```

4 Method Agreement Procedure:

The choice of 2 clusters is supported by 8 (27.59%) methods out of 29 (Elbow, Silhouette, Ch, Ratkowsky, Mcclain, SDindex, Mixture (VEV), Mixture (EEV)).

```
plot(n_clusters(data.frame(automobile)))
```



5 Two Step Clustering

Some people prefer using hierarchical methods first to obtain the number of clusters and their centroids. Then use these as inputs to run non-hierarchical analysis. This approach is called Two-Step Clustering.

```
# two step clustering in R, not required
library(prcr)
d <- pisaUSA15
m3 <- create_profiles_cluster(d,
                             broad_interest, enjoyment, instrumental_mot, self_efficac
                             n_profiles = 3)
summary(m3)
```

6 Clustering for Cases and Variables

Infrastructure construction includes various dimensions such as transportation, postal and telecommunications, water supply and power supply. “construction.xlsx” records the development level of 16 infrastructure construction aspects in 7 countries. Please answer the following questions:

1. Please use hierarchical methods to determine how many different types of infrastructure construction are included in these 16 indicators, and briefly explain the possible meanings of each type.
2. Calculate the average value of each type of indicator for each country in the previous question as the overall measure of infrastructure construction for that type (assuming this operation is reasonable). Use the k-means method to cluster the 7 countries based on their different overall indicators, distinguish between developed and developing countries, and determine which countries belong to each category.

6.1 Q1

```
library(factoextra)
library(showtext)
showtext_auto()
```

```

cons = import('/Users/zhourui/Desktop/Psychological Statistics/Homework/7/construction


## Successfully imported: 16 obs. of 8 variables

rownames(cons) = cons[,1]
cons = cons %>% select(-1)
attach(cons)

distance = dist(cons)
hc = hclust(distance)

plot(hc, hang = -1)

```



From the cluster dendrogram above, it's suggested to classify all 16 indexes into 3 grand categories. From indexes contained in each category, we can infer that the three dimensions are:

- Medical resources
- Education level
- Traffic capacity

6.2 Q2

```

grouplabel = cutree(hc,k = 3)

km_country = aggregate(cons,list(grouplabel),mean) %>%
  select(-1) %>% t() %>%
  kmeans(centers = 2)
country_label = km_country$cluster
country_label

```

```

## A B C D E F G
## 1 1 1 2 2 2 1

```

From the k -means method, we can see that the seven countries are classified into two grand categories, developed and developing countries, where the first group consists of D, E & F, and all else as the second group. However, clustering analysis doesn't directly tell us the attributes of two groups, therefore, we have to "label" the groups by checking the data.

Note that here D, E & F are labeled as group "2".

```

cons %>% t() %>% as.data.frame() %>% mutate(
  develop = country_label
) %>% group_by(develop) %>% summarise_all(mean) %>%
  stargazer::stargazer(header = FALSE,
                        title = 'Index Comparison Between Groups',
                        flip = T,summary = FALSE)

```

From the summarization table, apparently that among all the 16 indexes, group 1 outbeats group 2 (i.e., D, E & F). Therefore, we can say

- Developing countries: D, E, F;
- Developed countries: A, B, C, G.


```

fviz_cluster(km_country,
              data = aggregate(cons,list(grouplabel),mean) %>%
                select(-1) %>% t())

```

表 1: Index Comparison Between Groups

	1	2
develop	1	2
每十万人医疗机构数	361.035	214.980
每十万人医疗技术人员数	617.005	244.910
高速公路总里程	351.375	112.713
铁路总里程	374.023	119.360
每十万人教师数	158.267	78.060
每十万人学校数	63.805	26.087
每千人文盲数	115.995	63.923
高等教育入学率	26.145	6.663
在学儿童数量	129.545	62.687
航空里程数	244.132	103.263
航运总里程	290.225	135.253
铁路客运量	260.318	95.403
公路客运量	296.095	113.263
教育投入总费用	220.273	75.360
卫生投入总费用	489.022	229.947
每十万人医疗床位数	495.518	182.137



All done. Finally,

```
detach(cons)
```